

Introduction to LLMs in Python — Personal Study Guide

What this is. A study companion for your own learning, built from the skills in the DataCamp course you downloaded. Each section **explains the theory**, then gives you **original exercises**: for each one you get the task, a walk-through of the code you're handed, a hint, and the file to complete. **Solutions are in a separate document** (`study-guide-solutions.pdf`) — try each exercise first. Personal use only.

The runnable workbook is in `llm-exercises/`: **one self-contained file per exercise** under `exercises/` (imports at the top, `solve` to complete), answers under `solutions/`, and `check.py` (verifier).

1 · Setup & how to verify it works

```
cd llm-exercises
python3 -m venv .venv && source .venv/bin/activate # Windows: .venv\Scripts\activate
pip install -r requirements.txt
```

Confirm your environment before starting (exercises 4 and 9 need no internet):

```
python check.py 4 9 # expect two TODO lines – harness runs, work not done
python check.py --solutions 4 9 # expect: PASS PASS – reference solutions run
```

Then work through `exercises/` and re-run the verifier:

```
python check.py # checks all 11
python check.py 1 2 6 # only some
python check.py --show 4 11 # also PRINT what each solve() returned
python exercises/ex01_summarize.py # or run a single file standalone (prints its output)
```

Tip: `check.py` shows only PASS/FAIL by default. Add `--show` to also print each `solve()`'s return value, or run a file directly (`python solutions/ex01_summarize.py`) to see its output.

Each exercise reports one status:

Status	Meaning
PASS	Your <code>solve()</code> returned the right thing.
TODO	Not implemented yet (an <code>____</code> blank is still there).

Status	Meaning
FAIL	It ran but returned something wrong (the message says why).
SKIP	A model/dataset couldn't download (no internet), or deps missing.
ERR	Unexpected error (message shown).

Exercises other than 4 and 9 download a **small** model from the Hugging Face Hub on first use (internet + a few hundred MB). Everything runs on CPU except real training.

To complete an exercise, open its file, replace each `_____` blank (the comment above each line says what goes there), and re-run the verifier.

For a text exercise the verifier runs the file's own `SAMPLE`, so editing `SAMPLE` changes what you see (and `--show` prints *your* input's result); `PASS` means the result is **valid** — a real POSITIVE/NEGATIVE label, a non-empty translation, and so on — not that some fixed sentence matched.

2 · Using pre-trained models with pipelines

A **large language model (LLM)** is a deep-learning model — almost always a **Transformer** — with a huge number of parameters, trained on massive amounts of text. You rarely build one; you *use* one that already exists. Hugging Face's `transformers` library makes that a one-liner through the `pipeline()` function: you name a *task* and (optionally) a *model*, and it hands back a callable that does the whole job — tokenizing the input, running the model, and decoding the output.

Tasks fall into two broad families. **Understanding** tasks read text and return a structured result — a label or a span: sentiment analysis, classification, extractive question-answering. **Generation** tasks produce new text: text generation, summarization, translation.

For generation you usually set a few knobs. `max_length` caps how long the output can be; `truncation=True` trims over-long inputs; and `pad_token_id` tells the model which token to pad with — setting it to `tokenizer.eos_token_id` (the end-of-sequence token) silences the common padding warning. Beyond parameters, the strongest lever is the **prompt** itself: the words you feed in steer what comes out. **Translation** is just another task — either name it explicitly (`translation_en_to_es`) or pick a dedicated model such as the `Helsinki-NLP/opus-mt-*` family.

Exercises

Exercise 1 · Summarize with a pipeline

Task. Return a ≤ 40 -word summary of `text` using a `summarization` pipeline.

What the code does. The file imports `pipeline`, the one-line entry point to a pretrained model. `pipeline(task=..., model=...)` builds a ready-to-use summarizer around the DistilBART model; calling `summarizer(text, max_length=40)` runs it and returns a **list with one dict**, where `max_length=40` caps the summary length. The two blanks are the **task string** that tells `pipeline` which job to do, and the **dict key** that holds the produced summary.

Line by line:

- `from transformers import pipeline` — imports `pipeline`, the one-line helper that wraps a pretrained model into a ready-to-call object.
- `def solve(text):` — defines the function; `text` is the long passage you want to summarize.
- `summarizer = pipeline(task=____, model="sshleifer/distilbart-cnn-12-6")` — builds a **summarization** pipeline around the DistilBART model — this downloads the model from the Hugging Face Hub the first time. The blank is the *task name*.
- `out = summarizer(text, max_length=40)` — runs the model on `text`; `max_length=40` caps the summary length. It returns a **list with a single dict**, e.g.
`[{"summary_text": "..."}]`.
- `return out[0][____]` — takes the first (only) element and reads the summary string out of its dict. The blank is the *key* that holds it.

Why this model. Summarization is a *sequence-to-sequence* task — read a whole document, then write a shorter one — so it needs an **encoder-decoder** Transformer (see §3). `sshleifer/distilbart-cnn-12-6` is a **distilled BART** (a smaller, faster BART) fine-tuned on the CNN/DailyMail *news* dataset; the `12-6` are its 12 encoder and 6 decoder layers. That fine-tuning is exactly why it condenses text into fluent, news-style summaries out of the box — a plain GPT-style model would just keep writing, not summarize.

Hint. The task is named after what it does. The result is a one-element list of dicts, and the field name tells you which one holds the summary.

Complete this in `exercises/ex01_summarize.py`:

```
from transformers import pipeline

def solve(text):
    # Fill: the task name for summarization, and the output key.
    summarizer = pipeline(task=____, model="sshleifer/distilbart-cnn-12-6")
    out = summarizer(text, max_length=40)
    return out[0][____]
```

Exercise 2 · Generate text (control length + padding)

Task. Return ~30 **new** tokens of continuation from `distilgpt2`, with **no** warnings.

What the code does. `pipeline("text-generation", "distilgpt2")` wraps distilGPT-2; calling `gen(prompt, ...)` continues your prompt, and `truncation=True` trims an over-long prompt. The two blanks are `max_new_tokens` (how many **new** tokens to generate — the modern length control; `max_length` is the older one that counts the prompt too, and mixing both triggers a warning) and `pad_token_id` (which token id the model pads with — GPT-2 has no dedicated pad token, so you reuse the tokenizer's end-of-sequence token).

Line by line:

- `from transformers import pipeline` — imports `pipeline`, the one-line helper that wraps a pretrained model.
- `def solve(prompt):` — defines the function; `prompt` is the starting text the model will continue.
- `gen = pipeline(task="text-generation", model="distilgpt2")` — builds a **text-generation** pipeline around distilGPT-2 — downloaded from the Hub the first time.
- `out = gen(prompt, max_new_tokens=____, pad_token_id=gen.tokenizer.____, truncation=True)` — runs the model on `prompt`. `max_new_tokens` (blank) caps how many **new** tokens are generated — use this rather than `max_length` (the pipeline defaults `max_new_tokens`, so setting `max_length` too would warn). `pad_token_id` (blank) sets which token pads short outputs — GPT-2 has no pad token, so you reuse its end-of-sequence token, exposed on `gen.tokenizer`; `truncation=True` trims an over-long prompt. Returns a **list with one dict**.
- `return out[0]["generated_text"]` — reads the generated string (prompt + continuation) from the first result's dict, under the `generated_text` key.

Why this model. Text generation means predicting the next token over and over — which is what a **decoder-only** Transformer does (see §3). `distilgpt2` is a **distilled GPT-2** (~82M parameters, a lighter/faster GPT-2) small enough to download and run on a CPU. It is *not* fine-tuned for any task, so it simply **continues** whatever text you give it — which is precisely what lets you see raw next-token generation.

Hint. The task says how many *new* tokens. For padding, remember GPT-2 has no pad token, so it reuses its *end-of-sequence* token — look for that id on the tokenizer.

Complete this in `exercises/ex02_generate.py`:

```
from transformers import pipeline

def solve(prompt):
    gen = pipeline(task="text-generation", model="distilgpt2")
    # Fill: how many NEW tokens to generate (30), and the end-of-sequence token id
    # to pad with. (Use max_new_tokens, not max_length: the pipeline already
    # defaults max_new_tokens, so mixing in max_length triggers a warning.)
    out = gen(prompt, max_new_tokens=____,
              pad_token_id=gen.tokenizer.____, truncation=True)
    return out[0]["generated_text"]
```

Exercise 3 · Translate with a pipeline

Task. Return the English→Spanish translation of `text`.

What the code does. `pipeline(task=..., model="Helsinki-NLP/opus-mt-en-es")` builds a Marian EN→ES translator; calling it returns `[{"translation_text": ...}]`, which the code reads. The blank is the **task name**, which for translation encodes the direction.

Line by line:

- `from transformers import pipeline` — imports `pipeline`, the one-line helper that wraps a pretrained model.
- `def solve(text):` — defines the function; `text` is the English sentence to translate.
- `translator = pipeline(task=..., model="Helsinki-NLP/opus-mt-en-es")` — builds a **translation** pipeline around the Marian OPUS-MT English→Spanish model — downloaded from the Hub the first time. The blank is the *task name*.
- `return translator(text)[0]["translation_text"]` — runs it; the pipeline returns `[{"translation_text": ...}]`, and this reads the translated string from the first (only) dict.

Why this model. Translation is a *sequence-to-sequence* task — read a full sentence in one language, then write it in another — so, like summarization, it uses an **encoder-decoder** Transformer (see §3). `Helsinki-NLP/opus-mt-en-es` is a **MarianMT** model from the OPUS project, trained specifically on English→Spanish parallel text; its name encodes exactly that (`opus-mt-en-es` = OPUS machine translation, en→es). Being a small model dedicated to one language pair, it translates accurately and cheaply — a big general LLM could translate too, but this is the focused, efficient choice.

Hint. Translation task names encode the direction as `translation_<src>_to_<tgt>`; fill in the two language codes.

Complete this in `exercises/ex03_translate.py`:

```
from transformers import pipeline

def solve(text):
    # Fill: the task name for English -> Spanish translation.
    translator = pipeline(task=____,
                          model="Helsinki-NLP/opus-mt-en-es")
    return translator(text)[0]["translation_text"]
```

3 · The Transformer, in three shapes

Every modern language model is a **Transformer**, built from two kinds of block: an **encoder** (bidirectional attention — each token sees the whole input, so it *understands*)

and a **decoder** (causal attention — each token sees only what came before, which is what lets it *generate*). Keeping one or both gives three model shapes, each suited to different tasks. The figure below is the whole picture at a glance.

3 The Transformer, in three shapes




Every modern language model is a Transformer, stacked from two kinds of building blocks. Knowing how they differ explains what each model is good at — and why today's LLMs look the way they do.

The two building blocks: attention

1 The encoder: bidirectional attention

Every token sees every other token, left and right. Think once, it builds a rich, context-aware representation of each word.



- ✓ Understands input. It does not produce text.

2 The decoder: causal (masked) attention

Each token sees only the tokens before it, never the future. It can only predict the next token from what has been generated so far (autoregression).



- ✓ Generates text one token at a time. Peeking ahead would be cheating.

From these two blocks come three model shapes



1 Encoder-only (BERT family)

Just the encoder stack. Trained by hiding random words and predicting them from both sides of context, so it produces contextual embeddings without running text.

Best for understanding tasks: text classification, sentiment, NER, extractive QA...



2 Decoder-only (GPT family)

Just the decoder stack. Trained to predict the next token over enormous amounts of text. Being autoregressive, it generates.

Best for generation tasks: general-purpose LLMs (ChatGPT, Claude, Llama...)



3 Encoder-decoder (T5, BART, Marian)

Both stacks, joined by cross-attention: the decoder attends to its own past (causal) and to the encoder's representation of the input (bi-directional).

Best for sequence-to-sequence tasks: translation, summarization, data-to-text...



Check a model in code: `pipe.model.config.is_decoder` and `pipe.model.config.is_encoder_decoder` tell you which shape you loaded.



An important nuance

These families describe how models are *built* and what small, dedicated models are optimised for — not a hard limit on tasks.

A large **decoder-only** LLM (ChatGPT, Claude) happily summarizes and translates too: you put the document in the prompt, and it generates the summary as a continuation, using causal attention over the prompt. Enough scale plus instruction-tuning lets a single decoder stack do a lot.



Classic, efficient choice for focused small models: use the right shape for the job.

Shape	Attention	Trained to	Best at	Examples
 Encoder-only (BERT family)	Bidirectional 	Fill in masked words	Understand: classify, NER, extractive QA	BERT
 Decoder-only (GPT family)	Causal 	Predict the next token	Generate: general-purpose LLMs	GPT-2, ChatGPT, Claude, Llama
 Encoder-decoder (T5, BART, Marian)	Bi-directional + Cross-attention 	Map input → output	Seq-to-seq: translate, summarize, etc.	T5, BART, Marian, mT5

Exercises

Exercise 4 · Pick the right transformer family · *offline*

Task. Map a task key ("sentiment", "poem", "translation", "extractive-qa") to "encoder-only" / "decoder-only" / "encoder-decoder". *Offline*.

What the code does. Pure Python, no model. `solve(task)` looks a task up in a dictionary and returns which Transformer family fits it. The four blanks are the **family strings** for each task.

Hint. Sort each task by what it needs: to *understand* the input, to *generate* free text, or to *transform* one sequence into another.

Complete this in `exercises/ex04_family.py`:

```
def solve(task):
    # Fill each family: understanding -> "encoder-only";
    # free generation -> "decoder-only"; seq-to-seq -> "encoder-decoder".
    families = {
        "sentiment": ____,
        "poem": ____,
        "translation": ____,
        "extractive-qa": ____,
    }
    return families[task]
```

Exercise 5 · Classify sentiment

Task. Return "POSITIVE" / "NEGATIVE" for `text` using a fine-tuned sentiment pipeline.

What the code does. `pipeline(task=..., model="distilbert-base-uncased-finetuned-sst-2-english")` wraps a DistilBERT model fine-tuned on movie-review sentiment; calling it returns `[{"label": ..., "score": ...}]`. The blanks are the **task name** and the **key** holding the predicted label.

Line by line:

- `from transformers import pipeline` — imports the pipeline helper.
- `def solve(text):` — defines the function; `text` is the sentence to classify.
- `clf = pipeline(task=____, model="distilbert-base-uncased-finetuned-sst-2-english")` — builds a classification pipeline around a DistilBERT model fine-tuned for sentiment (downloaded the first time). The blank is the *task name*.
- `return clf(text)[0][____]` — runs it; the pipeline returns `[{"label": ..., "score": ...}]`, and this reads the predicted label. The blank is the *key*.

Why this model. `distilbert-base-uncased-finetuned-sst-2-english` is an **encoder-only** DistilBERT (§3) that has **already been fine-tuned** on SST-2 (a movie-review sentiment dataset) to output POSITIVE/NEGATIVE. Because the fine-tuning is done, the pipeline classifies out of the box; DistilBERT itself is a distilled (smaller, faster) BERT.

Hint. The task name is the common two-word phrase for judging sentiment; the result dict reports its prediction under an obvious key.

Complete this in `exercises/ex05_sentiment.py`:

```
from transformers import pipeline

def solve(text):
    # Fill: the task name for sentiment, and the output key holding the label.
    clf = pipeline(task=____,
                   model="distilbert-base-uncased-finetuned-sst-2-english")
    return clf(text)[0][____]
```

4 · From pipelines to Auto classes: datasets & tokenization

`pipeline()` is convenient but hides everything. To fine-tune — or simply to control the pieces — you drop down to the **Auto classes**: `AutoTokenizer` for the tokenizer and `AutoModelForSequenceClassification` (or `AutoModel`, ...) for the model, each loaded with `.from_pretrained(name)`.

It helps to picture the **LLM lifecycle**: **pre-training** on broad data teaches general language patterns; then optional **fine-tuning** on your own domain-specific data adapts the model to a task. That data usually comes from a dataset — `datasets.load_dataset(...)` — which you must turn into numbers.

That conversion is **tokenization**. A tokenizer splits text into tokens and maps them to `input_ids` (plus an `attention_mask`). The knobs you'll use constantly: `padding` (pad short sequences so a batch is rectangular), `truncation` + `max_length` (cap long ones), and `return_tensors="pt"` (return PyTorch tensors). Modern tokenizers are **subword**: rare words split into meaningful pieces, so nothing is ever truly "out of vocabulary". To tokenize a whole dataset efficiently, wrap the tokenizer in a function and apply it across the whole dataset, in batches, with `dataset.map(fn, ...)`.

Exercises

Exercise 6 · Load an Auto model + tokenizer

Task. Return `(tokenizer, model)`, where the model is an `AutoModelForSequenceClassification` with 2 labels.

What the code does. This is the control path (no pipeline).

`AutoTokenizer.from_pretrained(name)` downloads the tokenizer that matches the model. The blanks are **which Auto class** loads the model (it must add a classification head) and **how many labels** that head has.

Line by line:

- `from transformers import AutoTokenizer, AutoModelForSequenceClassification` — imports the tokenizer and model **Auto classes** — the control path used to fine-tune.
- `def solve():` — no input; it returns the two objects you need.
- `name = "distilbert-base-uncased"` — the base model id — plain DistilBERT, *not* fine-tuned for any task.
- `tokenizer = AutoTokenizer.from_pretrained(name)` — downloads the tokenizer that matches that model.
- `model = _____.from_pretrained(name, num_labels=_____)` — loads the model with a fresh classification head. Blanks = which Auto class to use and how many output labels.
- `return tokenizer, model` — hands both back — you need the pair to tokenize *and* to run/fine-tune the model.

Why this model. `distilbert-base-uncased` is the **plain**, *not*-fine-tuned DistilBERT (encoder-only). `AutoModelForSequenceClassification` bolts a **fresh, untrained** classification head on top — so its predictions are random until you fine-tune it. That's the whole point: this is the *starting point for training*, in contrast to `ex5`, which loaded the already-fine-tuned version.

How it classifies. DistilBERT turns the text into **contextual embeddings** — one vector per token. The vector of the special `[CLS]` token (the first one) is used as the **summary of the whole sentence**. The **classification head** (a small linear layer — in DistilBERT preceded by a pre-classifier + ReLU) maps that summary vector to one **logit** per class: raw, unbounded scores. To read them as probabilities you apply **softmax**; but to simply pick the predicted class you take the `argmax` of the logits directly — no softmax needed (that's what `ex11` does). One caveat specific to this exercise: the head is **freshly added and untrained**, so its logits are essentially random until you fine-tune the model (contrast with `ex5`, whose head is already trained).

Hint. Of the two imported classes, one adds a classification head — its name says so. How many classes does POSITIVE-vs-NEGATIVE have?

Complete this in `exercises/ex06_load_model.py`:

```
from transformers import AutoTokenizer, AutoModelForSequenceClassification

def solve():
    name = "distilbert-base-uncased"
    tokenizer = AutoTokenizer.from_pretrained(name)
    # Fill: which imported class loads a classification head, and how many labels.
    model = _____.from_pretrained(name, num_labels=_____)
    return tokenizer, model
```

Exercise 7 · Tokenize texts (padding, truncation, tensors)

Task. Return padded, truncated (max 64) PyTorch tensors for a list of `texts`.

What the code does. The code calls the tokenizer on a list of texts. `padding=True` pads sequences to equal length so they form a rectangular batch; `truncation=True` cuts overly long ones. The blanks are the **tensor format** to return and the **max length** cap.

Line by line:

- `from transformers import AutoTokenizer` — imports the tokenizer class (used by the demo).
- `def solve(tokenizer, texts):` — receives a ready tokenizer and a list of texts.
- `return tokenizer(texts, return_tensors=____, padding=True, truncation=True, max_length=____)` — tokenizes the batch: `padding=True` pads all to equal length (a rectangular batch), `truncation=True` cuts long ones. Blanks = the tensor format and the length cap.

`input_ids` & `attention_mask`. The tokenizer returns a small dictionary of tensors:

- `input_ids` — the text turned into numbers: each token maps to an integer id from the model's vocabulary. For DistilBERT, `101` and `102` are the special `[CLS]` / `[SEP]` markers that wrap every sentence, and `0` is `[PAD]`.
- `attention_mask` — a parallel array, same shape, with **1 = real token** and **0 = padding**. *How it protects the result:* inside the Transformer, attention makes each token look at the others through weights produced by a **softmax**. At every position where the mask is 0, the attention score is set to $-\infty$ **before the softmax**, and since $e^{(-\infty)} = 0$ those slots get weight **0** — so padding contributes nothing and real tokens attend only to real tokens. That's why you can pad as much as you like without changing the output.
- **Not the same as the causal mask** (§3): `attention_mask` says *which tokens are real*; the causal mask (decoders only) says *each token may see only earlier ones*. The model adds the causal one itself — the tokenizer gives you the padding one.
- You always pass both to the model together: `model(**enc)` (i.e. `model(input_ids=..., attention_mask=...)`).

Hint. PyTorch tensors have a two-letter code; the length cap is the number named in the task.

Complete this in `exercises/ex07_tokenize.py`:

```
from transformers import AutoTokenizer

def solve(tokenizer, texts):
    # Fill: return PyTorch tensors ("pt"), and cap the length at 64 tokens.
    return tokenizer(texts, return_tensors=____, padding=True,
                    truncation=True, max_length=____)
```

Exercise 8 · Tokenize a whole dataset with `.map`

Task. Tokenize a dataset row-by-row with `.map`, processing it in batches.

What the code does. `load_dataset("imdb", split="train[:1%]")` loads 1% of the train split; the inner `tok` function tokenizes the `text` column; `.map` applies it to every row. The blank controls whether `.map` passes rows **one at a time or in batches**.

The point. The goal of this exercise isn't to train the model — it's to **prepare the data**. The raw dataset held only `text` and its `label`; running it through `.map` with the tokenizer produces a new dataset that keeps those and adds `input_ids` and `attention_mask` — the format Transformers expect for fine-tuning. In effect you've turned a **human-readable** dataset into one a **neural network can process**: `input_ids` and `attention_mask` are the model's **inputs**, `label` is the **target** it must learn to predict, and `text` just rides along. (One technicality: right after `.map` the ids are still Python lists; the Trainer / data collator turns them into tensors at training time.)

Line by line:

- `from transformers import AutoTokenizer` — the tokenizer class (used by the demo).
- `from datasets import load_dataset` — the loader for ready-made datasets.
- `def solve(tokenizer):` — receives a tokenizer to apply.
- `data = load_dataset("imdb", split="train[:1%]")` — loads 1% of IMDB's train split — a small set of labelled movie reviews.
- `def tok(batch): ...` — a helper that tokenizes the `text` column of a batch of rows.
- `return data.map(tok, batched=___)` — applies `tok` across the whole dataset. Blank = process rows **in batches** (much faster).

Batched tokenization with `.map`. `data.map(tok, ...)` applies the tokenizer across the **whole dataset** instead of one sentence at a time, and returns a **new dataset** that keeps the original columns (`text`, `label`) and **adds** `input_ids` and `attention_mask` — the columns the model consumes during fine-tuning. Three details worth knowing:

- **Batch size:** processing *in batches* doesn't mean *all at once* — by default `tok` receives **1,000 rows** per call.
- **Where the speed comes from:** the big win is that **fast (Rust) tokenizers** process a whole batch in parallel; with a slow (pure-Python) tokenizer the gain is much smaller.
- **A padding gotcha:** `padding=True` here pads **within each batch**, so different batches can end up with different lengths. For training you often defer padding to a *data collator* (dynamic per-batch padding) or use `padding="max_length"` for a uniform size.

Hint. You want `.map` to process rows in batches rather than one at a time.

Complete this in `exercises/ex08_tokenize_dataset.py`:

```

from transformers import AutoTokenizer
from datasets import load_dataset

def solve(tokenizer):
    data = load_dataset("imdb", split="train[:1%]")

    def tok(batch):
        return tokenizer(batch["text"], padding=True, truncation=True, max_length=128)

    # Fill: apply tok across the dataset in batches.
    return data.map(tok, batched=____)

```

5 · Fine-tuning through training

With a model and a tokenized dataset in hand, fine-tuning comes down to three objects.

TrainingArguments holds the hyperparameters — where to save (`output_dir`), when to evaluate (`eval_strategy="epoch"`), how long (`num_train_epochs`), how fast (`learning_rate`), the batch sizes, and regularization (`weight_decay`). **Trainer** ties everything together — the model, those args, the train and eval datasets, and the tokenizer — and `trainer.train()` runs the loop.

Once trained, using the model directly means doing by hand what the pipeline did for you: tokenize the text, run it under `torch.no_grad()` (no gradients needed for inference), read the raw **logits**, take the `argmax` for the predicted class id, and map that id to a label via `model.config.id2label`. Finally, `model.save_pretrained(dir)` and `tokenizer.save_pretrained(dir)` persist your work; `from_pretrained(dir)` loads it back.

Heads-up: in recent `transformers`, `TrainingArguments` and `Trainer` need the `accelerate` package (it's in `requirements.txt`). If you see "the Trainer ... requires accelerate", run `pip install -r requirements.txt` (or `pip install "accelerate>=1.1.0"`).

Exercises

Exercise 9 · Set up TrainingArguments · offline

Task. Return `TrainingArguments`: output `./out`, evaluate each epoch, 3 epochs, lr 2e-5, train/eval batch 8, weight decay 0.01. *Offline*.

What the code does. This constructs the hyperparameter bundle for training.

`output_dir` is where checkpoints go; each blank is one hyperparameter — when to evaluate, how many epochs, the learning rate, the train batch size, and weight decay — with its **target value written in the comment beside it**.

The point. This exercise doesn't train anything — it **defines the rules of training**.

`TrainingArguments` is a configuration object the `Trainer` reads later to know how many passes to make over the data (**epochs**), how big each weight update is (**learning rate**), how many examples to process at once (**batch size**), how much to regularise (**weight decay**), when to **evaluate**, and where to **save checkpoints** (`output_dir`). It's the *flight plan* the fine-tuning run will follow — it has dozens of options, and here you set the essential ones.

Line by line:

- `from transformers import TrainingArguments` — imports the hyperparameter container.
- `def solve():` — returns the configuration object (no training happens here).
- `return TrainingArguments(output_dir="./out", eval_strategy=____, num_train_epochs=____, ...)` — builds the training config; `output_dir` is where checkpoints go and each blank is one hyperparameter (eval schedule, epochs, learning rate, train batch size, weight decay) — the target value is written in the comment beside each blank.

Key hyperparameters. Three settings are worth understanding, not just filling in:

- `num_train_epochs` (**epochs**) — one epoch is one full pass over the whole training set, so `3` means the model sees every example three times. Too few and it *underfits* (hasn't learned enough); too many and it starts to *memorise* the training data (**overfits**).
- `learning_rate` — the **size of each step** when the weights are updated. Big steps learn fast but can overshoot and diverge; tiny steps are stable but slow. `2e-5` (`0.00002`) is a deliberately small value typical of fine-tuning: the pretrained weights are already good, so you only want to *nudge* them, not overwrite them.
- `weight_decay` — a **regularisation** term that, each step, gently pulls every weight toward zero (it penalises large weights). The penalty grows with the **square** of each weight (L2), so a few big weights cost far more than many small ones — e.g. two weights of 1 cost $1+1 = 2$, but one weight of 2 costs $2 \times 2 = 4$. The model therefore prefers to **spread** the signal across many features rather than lean hard on one, which **reduces overfitting**. It's a lever, though: too strong and the model is over-constrained and *underfits*. `0.01` is a mild, standard value. (Technical note: this is L2 regularisation; the `Trainer`'s AdamW optimiser applies it as *decoupled* weight decay — directly on the weights — but the intuition is the same.)

Hint. Each value you need is written in the comment next to its blank.

Complete this in `exercises/ex09_training_args.py`:

```

from transformers import TrainingArguments

def solve():
    # Fill the values described in the task.
    return TrainingArguments(
        output_dir="./out",
        eval_strategy=____,          # evaluate at the end of each epoch
        num_train_epochs=____,      # 3
        learning_rate=____,         # 2e-5
        per_device_train_batch_size=____, # 8
        per_device_eval_batch_size=8,
        weight_decay=____,          # 0.01
    )

```

Exercise 10 · Wire up the Trainer

Task. Return a `Trainer` built from the pieces (don't call `.train()`).

What the code does. In a single call you build the `Trainer` — the object that will run the fine-tuning loop — handing it the pieces from the previous exercises. `solve()` only *constructs* it; no training happens here.

The point. This is the moment **every piece you built connects into one pipeline**. There's no need to wire the arguments one by one to learn them — `model` (ex6), `TrainingArguments` (ex9), the train/eval datasets (ex8) and the `tokenizer` (ex6/7) were each covered already. The new idea is that the `Trainer` **orchestrates the whole fine-tuning**: hand it those pieces and it knows how to run the training loop.

Is the model being trained in this exercise? It depends what you run:

- `solve()` — the part you complete — only **builds** the `Trainer` object. **No training.**
- **Running the file** (`python exercises/ex10_build_trainer.py`) builds it *and* then calls `trainer.train()` on a tiny dataset for two steps — **yes, a real mini-train** you can watch (a loss is printed).
- `python check.py 10` only verifies the wiring. **No training.**

So the *exercise itself* just assembles the `Trainer`; the bundled *demo* adds a two-step train so the loop isn't abstract. And what `trainer.train()` actually does, each step, is: pull a **batch** from the `train_dataset` → run it **through the model** → compute the **loss** (how wrong it is) → **backpropagate** to get the gradients (which way each weight should move) → **update the weights** a little (step size = the learning rate) → **evaluate** when `eval_strategy` says to → **save checkpoints** to `output_dir`.

Hint. Each `Trainer` parameter takes the argument of the same name that `solve` received.

Complete this in `exercises/ex10_build_trainer.py`:

```

from transformers import (AutoTokenizer, AutoModelForSequenceClassification,
                          Trainer, TrainingArguments)
from datasets import Dataset

def solve(model, args, train_ds, eval_ds, tokenizer):
    # Fill: pass the model, the args, and the two datasets into the Trainer.
    return Trainer(
        model=model,
        args=args,
        train_dataset=train_ds,
        eval_dataset=eval_ds,
        tokenizer=tokenizer,
    )

```

Exercise 11 · Use the fine-tuned model (from logits)

Task. Return "POSITIVE" / "NEGATIVE" for `text`, computed **from raw logits**.

The key idea. The goal isn't to learn `torch.argmax()` — it's to understand that a classification model **doesn't return a label**; it returns **logits** (raw numeric scores), and *you* have to turn those into the final class. This exercise does by hand what `pipeline()` was doing for you automatically.

Where we are. This is the last step of *inference*:

```
Text -> Tokenizer -> Model -> Logits -> pick the best class -> "POSITIVE"
```

What are logits? The model doesn't hand back `POSITIVE`; it returns a vector of scores, e.g.

```
tensor([[ -1.25,  3.82]])
```

one score per class (index `0` -> NEGATIVE, index `1` -> POSITIVE). These are **not** probabilities — they're raw scores. (You'd apply *softmax* to turn them into probabilities, but you don't need to just to pick the winner.)

How we get the prediction. `torch.argmax(logits, dim=1)` returns the *position* of the largest value. For `[-1.25, 3.82]` that is `1`, because `3.82` is the highest.

Why `dim=1`? Logits have shape `[batch, classes]` — one row per sentence, one column per class. Reducing over `dim=1` (the class axis) picks the top class **for each sentence** in the batch.

How we get the label. The model carries a dictionary `model.config.id2label`, e.g. `{0: "NEGATIVE", 1: "POSITIVE"}`, so `model.config.id2label[pred]` turns the class id into a human-readable label.

Why `torch.no_grad()` ? At inference we don't want to compute gradients or update weights — we only read a prediction. Wrapping the forward pass in `with torch.no_grad():` makes it faster and use less memory.

Why the model matters. This is the same SST-2 fine-tuned DistilBERT as ex5. *Because* it's fine-tuned, its logits are meaningful and `argmax` gives a real POSITIVE/NEGATIVE; the untrained head from ex6 would give noise.

The full flow:

```
Text
|
v
Tokenizer -> input_ids
|
v
fine-tuned Model
|
v
Logits          e.g. [-1.25, 3.82]
|
v   argmax()    -> 1
id2label[1]
|
v
"POSITIVE"
```

What you really learn. This exercise reveals what was hidden inside `pipeline()`: a classification model returns a vector of **logits**, and making a prediction means selecting the highest-scoring class (`argmax`) and mapping its id to a label (`id2label`). That is what actually happens during inference with a Transformer.

Hint. Logits are shaped `[batch, classes]`, so reduce over the class axis; then match the capitalisation the task expects.

Complete this in `exercises/ex11_predict.py`:

```
import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification

def solve(text):
    name = "distilbert-base-uncased-finetuned-sst-2-english"
    tok = AutoTokenizer.from_pretrained(name)
    model = AutoModelForSequenceClassification.from_pretrained(name)
    inp = tok(text, return_tensors="pt", padding=True, truncation=True)
    with torch.no_grad():
        logits = model(**inp).logits
    # Fill: the axis to argmax over (1), and the string method to upper-case.
    pred = int(torch.argmax(logits, dim=____))
    return model.config.id2label[pred].____()
```

6 · Fine-tuning approaches & N-shot prompting

Full training isn't the only option. **Full fine-tuning** updates every weight — most powerful, most expensive. **Partial fine-tuning** freezes most layers and trains only a few — cheaper and less prone to overfitting. **Transfer learning** is the umbrella idea: take a model trained on one task and adapt it to a related one.

Often you don't need to train at all. **N-shot prompting** teaches the model *inside the prompt*, by example: **zero-shot** gives no example (just the instruction), **one-shot** gives a single worked example, and **few-shot** gives several. More examples generally sharpen the output format and accuracy — at the cost of a longer prompt.

7 · Where to look things up

Start here (official, free):

- **Hugging Face LLM/NLP Course** — the best structured, hands-on intro: <https://huggingface.co/learn>
- **Transformers docs (home)** — <https://huggingface.co/docs/transformers>

API reference for what you used in these exercises:

- `pipeline()` and the full task list — https://huggingface.co/docs/transformers/main_classes/pipelines
- **Auto classes** (`AutoTokenizer`, `AutoModelFor...`) — https://huggingface.co/docs/transformers/model_doc/auto
- **Tokenizer** (padding / truncation / `return_tensors`) — https://huggingface.co/docs/transformers/main_classes/tokenizer
- **Trainer & TrainingArguments** — https://huggingface.co/docs/transformers/main_classes/trainer
- **Datasets** (`load_dataset`, `.map`) — <https://huggingface.co/docs/datasets>
- **Task guides** (summarization, translation, classification, QA...) — <https://huggingface.co/docs/transformers/tasks/summarization>

Finding a model:

- **The Model Hub** — filter by task on the left: <https://huggingface.co/models>
- **Read the model card** — what architecture it is, its intended task and languages, and the exact `model=` id to copy.

Background / the "why":

- PyTorch docs — <https://pytorch.org/docs/stable/index.html>
- Papers: *Attention Is All You Need* (Transformer, 2017), *BERT* (2018), *GPT-2* (2019) — searchable on <https://arxiv.org>

A skill worth practising, not just bookmarking:

- In Python, `help(pipeline)` (or `TrainingArguments?` in a notebook) prints the docstring for any object — the fastest reference of all.
 - Every argument you pass (`max_length`, `eval_strategy`, ...) has a one-line explanation in the class's docstring / API page — look it up there instead of guessing.
-

8 • Self-check

You can now: run a task with `pipeline()`; control generation length/padding; translate; tell encoder-only / decoder-only / encoder-decoder apart and verify in code; move from `pipeline` to Auto classes; tokenize (with padding/truncation) and `.map` a dataset; set up `TrainingArguments` + `Trainer`; run inference from logits; and prompt zero/one/few-shot. That's the whole arc of the two chapters.