

Introduction to LLMs in Python — Personal Study Guide

What this is. A study companion for your own learning, built from the skills in the DataCamp course you downloaded. Each section **explains the theory**, then gives you **original exercises**: for each one you get the task, a walk-through of the code you're handed, a hint, and the file to complete. **Solutions are in a separate document** (`study-guide-solutions.pdf`) — try each exercise first. Personal use only.

The runnable workbook is in `llm-exercises/`: **one self-contained file per exercise** under `exercises/` (imports at the top, `solve` to complete), answers under `solutions/`, and `check.py` (verifier).

1 · Setup & how to verify it works

```
cd llm-exercises
python3 -m venv .venv && source .venv/bin/activate # Windows: .venv\Scripts\activate
pip install -r requirements.txt
```

Confirm your environment before starting (exercises 4 and 9 need no internet):

```
python check.py 4 9 # expect two TODO lines – harness runs, work not done
python check.py --solutions 4 9 # expect: PASS PASS – reference solutions run
```

Then work through `exercises/` and re-run the verifier:

```
python check.py # checks all 11
python check.py 1 2 6 # only some
python check.py --show 4 11 # also PRINT what each solve() returned
python exercises/ex01_summarize.py # or run a single file standalone (prints its output)
```

Tip: `check.py` shows only PASS/FAIL by default. Add `--show` to also print each `solve()`'s return value, or run a file directly (`python solutions/ex01_summarize.py`) to see its output.

Each exercise reports one status:

Status	Meaning
PASS	Your <code>solve()</code> returned the right thing.
TODO	Not implemented yet (an <code>____</code> blank is still there).

Status	Meaning
FAIL	It ran but returned something wrong (the message says why).
SKIP	A model/dataset couldn't download (no internet), or deps missing.
ERR	Unexpected error (message shown).

Exercises other than 4 and 9 download a **small** model from the Hugging Face Hub on first use (internet + a few hundred MB). Everything runs on CPU except real training.

To complete an exercise, open its file, replace each `_____` blank (the comment above each line says what goes there), and re-run the verifier.

For a text exercise the verifier runs the file's own `SAMPLE`, so editing `SAMPLE` changes what you see (and `--show` prints *your* input's result); `PASS` means the result is **valid** — a real POSITIVE/NEGATIVE label, a non-empty translation, and so on — not that some fixed sentence matched.

2 · Using pre-trained models with pipelines

A **large language model (LLM)** is a deep-learning model — almost always a **Transformer** — with a huge number of parameters, trained on massive amounts of text. You rarely build one; you *use* one that already exists. Hugging Face's `transformers` library makes that a one-liner through the `pipeline()` function: you name a *task* and (optionally) a *model*, and it hands back a callable that does the whole job — tokenizing the input, running the model, and decoding the output.

Tasks fall into two broad families. **Understanding** tasks read text and return a structured result — a label or a span: sentiment analysis, classification, extractive question-answering. **Generation** tasks produce new text: text generation, summarization, translation.

For generation you usually set a few knobs. `max_length` caps how long the output can be; `truncation=True` trims over-long inputs; and `pad_token_id` tells the model which token to pad with — setting it to `tokenizer.eos_token_id` (the end-of-sequence token) silences the common padding warning. Beyond parameters, the strongest lever is the **prompt** itself: the words you feed in steer what comes out. **Translation** is just another task — either name it explicitly (`translation_en_to_es`) or pick a dedicated model such as the `Helsinki-NLP/opus-mt-*` family.

Exercises

Exercise 1 · Summarize with a pipeline

Task. Return a ≤ 40 -word summary of `text` using a `summarization` pipeline.

What the code does. The file imports `pipeline`, the one-line entry point to a pretrained model. `pipeline(task=..., model=...)` builds a ready-to-use summarizer around the DistilBART model; calling `summarizer(text, max_length=40)` runs it and returns a **list with one dict**, where `max_length=40` caps the summary length. The two blanks are the **task string** that tells `pipeline` which job to do, and the **dict key** that holds the produced summary.

Line by line:

- `from transformers import pipeline` — imports `pipeline`, the one-line helper that wraps a pretrained model into a ready-to-call object.
- `def solve(text):` — defines the function; `text` is the long passage you want to summarize.
- `summarizer = pipeline(task=____, model="sshleifer/distilbart-cnn-12-6")` — builds a **summarization** pipeline around the DistilBART model — this downloads the model from the Hugging Face Hub the first time. The blank is the *task name*.
- `out = summarizer(text, max_length=40)` — runs the model on `text`; `max_length=40` caps the summary length. It returns a **list with a single dict**, e.g.
`[{"summary_text": "..."}]`.
- `return out[0][____]` — takes the first (only) element and reads the summary string out of its dict. The blank is the *key* that holds it.

Why this model. Summarization is a *sequence-to-sequence* task — read a whole document, then write a shorter one — so it needs an **encoder-decoder** Transformer (see §3). `sshleifer/distilbart-cnn-12-6` is a **distilled BART** (a smaller, faster BART) fine-tuned on the CNN/DailyMail *news* dataset; the `12-6` are its 12 encoder and 6 decoder layers. That fine-tuning is exactly why it condenses text into fluent, news-style summaries out of the box — a plain GPT-style model would just keep writing, not summarize.

Hint. The task is named after what it does. The result is a one-element list of dicts, and the field name tells you which one holds the summary.

Complete this in `exercises/ex01_summarize.py`:

```
from transformers import pipeline

def solve(text):
    # Fill: the task name for summarization, and the output key.
    summarizer = pipeline(task=____, model="sshleifer/distilbart-cnn-12-6")
    out = summarizer(text, max_length=40)
    return out[0][____]
```

Exercise 2 · Generate text (control length + padding)

Task. Return ~30 **new** tokens of continuation from `distilgpt2`, with **no** warnings.

What the code does. `pipeline("text-generation", "distilgpt2")` wraps distilGPT-2; calling `gen(prompt, ...)` continues your prompt, and `truncation=True` trims an over-long prompt. The two blanks are `max_new_tokens` (how many **new** tokens to generate — the modern length control; `max_length` is the older one that counts the prompt too, and mixing both triggers a warning) and `pad_token_id` (which token id the model pads with — GPT-2 has no dedicated pad token, so you reuse the tokenizer's end-of-sequence token).

Line by line:

- `from transformers import pipeline` — imports `pipeline`, the one-line helper that wraps a pretrained model.
- `def solve(prompt):` — defines the function; `prompt` is the starting text the model will continue.
- `gen = pipeline(task="text-generation", model="distilgpt2")` — builds a **text-generation** pipeline around distilGPT-2 — downloaded from the Hub the first time.
- `out = gen(prompt, max_new_tokens=____, pad_token_id=gen.tokenizer.____, truncation=True)` — runs the model on `prompt`. `max_new_tokens` (blank) caps how many **new** tokens are generated — use this rather than `max_length` (the pipeline defaults `max_new_tokens`, so setting `max_length` too would warn). `pad_token_id` (blank) sets which token pads short outputs — GPT-2 has no pad token, so you reuse its end-of-sequence token, exposed on `gen.tokenizer`; `truncation=True` trims an over-long prompt. Returns a **list with one dict**.
- `return out[0]["generated_text"]` — reads the generated string (prompt + continuation) from the first result's dict, under the `generated_text` key.

Why this model. Text generation means predicting the next token over and over — which is what a **decoder-only** Transformer does (see §3). `distilgpt2` is a **distilled GPT-2** (~82M parameters, a lighter/faster GPT-2) small enough to download and run on a CPU. It is *not* fine-tuned for any task, so it simply **continues** whatever text you give it — which is precisely what lets you see raw next-token generation.

Hint. The task says how many *new* tokens. For padding, remember GPT-2 has no pad token, so it reuses its *end-of-sequence* token — look for that id on the tokenizer.

Complete this in `exercises/ex02_generate.py`:

```
from transformers import pipeline

def solve(prompt):
    gen = pipeline(task="text-generation", model="distilgpt2")
    # Fill: how many NEW tokens to generate (30), and the end-of-sequence token id
    # to pad with. (Use max_new_tokens, not max_length: the pipeline already
    # defaults max_new_tokens, so mixing in max_length triggers a warning.)
    out = gen(prompt, max_new_tokens=____,
              pad_token_id=gen.tokenizer.____, truncation=True)
    return out[0]["generated_text"]
```

Exercise 3 · Translate with a pipeline

Task. Return the English→Spanish translation of `text`.

What the code does. `pipeline(task=..., model="Helsinki-NLP/opus-mt-en-es")` builds a Marian EN→ES translator; calling it returns `[{"translation_text": ...}]`, which the code reads. The blank is the **task name**, which for translation encodes the direction.

Line by line:

- `from transformers import pipeline` — imports `pipeline`, the one-line helper that wraps a pretrained model.
- `def solve(text):` — defines the function; `text` is the English sentence to translate.
- `translator = pipeline(task=____, model="Helsinki-NLP/opus-mt-en-es")` — builds a **translation** pipeline around the Marian OPUS-MT English→Spanish model — downloaded from the Hub the first time. The blank is the *task name*.
- `return translator(text)[0]["translation_text"]` — runs it; the pipeline returns `[{"translation_text": ...}]`, and this reads the translated string from the first (only) dict.

Why this model. Translation is a *sequence-to-sequence* task — read a full sentence in one language, then write it in another — so, like summarization, it uses an **encoder-decoder** Transformer (see §3). `Helsinki-NLP/opus-mt-en-es` is a **MarianMT** model from the OPUS project, trained specifically on English→Spanish parallel text; its name encodes exactly that (`opus-mt-en-es` = OPUS machine translation, en→es). Being a small model dedicated to one language pair, it translates accurately and cheaply — a big general LLM could translate too, but this is the focused, efficient choice.

Hint. Translation task names encode the direction as `translation_<src>_to_<tgt>`; fill in the two language codes.

Complete this in `exercises/ex03_translate.py`:

```
from transformers import pipeline

def solve(text):
    # Fill: the task name for English -> Spanish translation.
    translator = pipeline(task=____,
                          model="Helsinki-NLP/opus-mt-en-es")
    return translator(text)[0]["translation_text"]
```

3 · The Transformer, in three shapes

Every modern language model is a **Transformer**, stacked from two kinds of building block. Knowing how they differ explains what each model is good at — and why today's LLMs look the way they do.

At the heart of both blocks is **attention**: a mechanism that lets each token look at other tokens and pull in the information it needs to represent itself in context. The two blocks differ in *which* tokens a position is allowed to look at.

- **The encoder** uses **bidirectional** attention: every token sees every other token, left *and* right. Taking in the whole input at once, it builds a rich, context-aware representation of each word — it can tell that "bank" means the river kind from the words around it. An encoder does not produce text; it *understands* input.
- **The decoder** uses **causal** (masked) attention: each token sees only the tokens *before* it, never the future. That restriction is exactly what makes **generation** possible — the model predicts the next token from the ones so far, appends it, and repeats (this is *autoregression*). Letting it peek ahead would be cheating, since at generation time the future tokens don't exist yet.

From these two blocks come three model shapes:

Encoder-only (BERT family). Just the encoder stack. It is trained by hiding random words and predicting them from *both* sides of context, so it produces contextual embeddings but no running text. Ideal for *understanding* tasks: text classification, sentiment, named-entity recognition, extractive QA (selecting the answer span). Exercises 4 and 5 use one.

Decoder-only (GPT family). Just the decoder stack, trained on one simple objective — predict the next token over enormous amounts of text. Being autoregressive, it *generates*. This is the shape behind GPT, ChatGPT, Claude and Llama; because it scales so well on a single objective, it has become the default for large LLMs. Exercise 2 uses distilGPT-2.

Encoder-decoder (T5, BART, Marian). Both stacks, joined by **cross-attention**: the encoder reads the input bidirectionally, and the decoder generates the output while attending both to its own previous tokens (causal) *and* to the encoder's representation of the input (cross-attention). That "read all of the input, then write the output" structure is a natural fit for **sequence-to-sequence** tasks — which is why exercises 1 (DistilBART, summarization) and 3 (Marian, translation) use this shape.

Checking a model in code: `pipe.model.config.is_decoder` and `pipe.model.config.is_encoder_decoder` tell you which shape you loaded.

An important nuance. These families describe how models are *built* and what small, dedicated models are optimised for — not a hard limit on tasks. A large **decoder-only** LLM (ChatGPT, Claude) happily summarizes and translates too: you put the document in

the prompt and it generates the summary as a continuation, doing the "understanding" with causal attention over the prompt. Enough scale plus instruction-tuning lets a single decoder stack cover everything through prompting. So "summarization → encoder-decoder" is the *classic, efficient* choice for a focused small model (like the ones in these exercises), while a general chat LLM does it all with one decoder.

Shape	Attention	Trained to	Best at	Examples
Encoder-only	bidirectional	fill in masked words	understand: classify, NER, extractive QA	BERT
Decoder-only	causal	predict the next token	generate; general-purpose LLMs	GPT-2, ChatGPT, Claude
Encoder-decoder	bi + cross	map input → output	seq-to-seq: translate, summarize	T5, BART, Marian

Exercises

Exercise 4 · Pick the right transformer family · *offline*

Task. Map a task key ("sentiment", "poem", "translation", "extractive-qa") to "encoder-only" / "decoder-only" / "encoder-decoder" . *Offline*.

What the code does. Pure Python, no model. `solve(task)` looks a task up in a dictionary and returns which Transformer family fits it. The four blanks are the **family strings** for each task.

Hint. Sort each task by what it needs: to *understand* the input, to *generate* free text, or to *transform* one sequence into another.

Complete this in `exercises/ex04_family.py` :

```
def solve(task):
    # Fill each family: understanding -> "encoder-only";
    # free generation -> "decoder-only"; seq-to-seq -> "encoder-decoder".
    families = {
        "sentiment": ____,
        "poem": ____,
        "translation": ____,
        "extractive-qa": ____,
    }
    return families[task]
```

Exercise 5 · Classify sentiment

Task. Return "POSITIVE" / "NEGATIVE" for `text` using a fine-tuned sentiment pipeline.

What the code does. `pipeline(task=..., model="distilbert-base-uncased-finetuned-sst-2-english")` wraps a DistilBERT model fine-tuned on movie-review sentiment; calling

it returns `[{"label": ..., "score": ...}]`. The blanks are the **task name** and the **key** holding the predicted label.

Line by line:

- `from transformers import pipeline` — imports the pipeline helper.
- `def solve(text):` — defines the function; `text` is the sentence to classify.
- `clf = pipeline(task=____, model="distilbert-base-uncased-finetuned-sst-2-english")` — builds a classification pipeline around a DistilBERT model fine-tuned for sentiment (downloaded the first time). The blank is the *task name*.
- `return clf(text)[0][____]` — runs it; the pipeline returns `[{"label": ..., "score": ...}]`, and this reads the predicted label. The blank is the *key*.

Why this model. `distilbert-base-uncased-finetuned-sst-2-english` is an **encoder-only** DistilBERT (§3) that has **already been fine-tuned** on SST-2 (a movie-review sentiment dataset) to output POSITIVE/NEGATIVE. Because the fine-tuning is done, the pipeline classifies out of the box; DistilBERT itself is a distilled (smaller, faster) BERT.

Hint. The task name is the common two-word phrase for judging sentiment; the result dict reports its prediction under an obvious key.

Complete this in `exercises/ex05_sentiment.py`:

```
from transformers import pipeline

def solve(text):
    # Fill: the task name for sentiment, and the output key holding the label.
    clf = pipeline(task=____,
                   model="distilbert-base-uncased-finetuned-sst-2-english")
    return clf(text)[0][____]
```

4 · From pipelines to Auto classes: datasets & tokenization

`pipeline()` is convenient but hides everything. To fine-tune — or simply to control the pieces — you drop down to the **Auto classes**: `AutoTokenizer` for the tokenizer and `AutoModelForSequenceClassification` (or `AutoModel`, ...) for the model, each loaded with `.from_pretrained(name)`.

It helps to picture the **LLM lifecycle**: **pre-training** on broad data teaches general language patterns; then optional **fine-tuning** on your own domain-specific data adapts the model to a task. That data usually comes from a dataset — `datasets.load_dataset(...)` — which you must turn into numbers.

That conversion is **tokenization**. A tokenizer splits text into tokens and maps them to `input_ids` (plus an `attention_mask`). The knobs you'll use constantly: `padding` (pad short

sequences so a batch is rectangular), `truncation + max_length` (cap long ones), and `return_tensors="pt"` (return PyTorch tensors). Modern tokenizers are **subword**: rare words split into meaningful pieces, so nothing is ever truly "out of vocabulary". To tokenize a whole dataset efficiently, wrap the tokenizer in a function and apply it across the whole dataset, in batches, with `dataset.map(fn, ...)`.

Exercises

Exercise 6 · Load an Auto model + tokenizer

Task. Return `(tokenizer, model)`, where the model is an `AutoModelForSequenceClassification` with 2 labels.

What the code does. This is the control path (no pipeline). `AutoTokenizer.from_pretrained(name)` downloads the tokenizer that matches the model. The blanks are **which Auto class** loads the model (it must add a classification head) and **how many labels** that head has.

Line by line:

- `from transformers import AutoTokenizer, AutoModelForSequenceClassification` — imports the tokenizer and model **Auto classes** — the control path used to fine-tune.
- `def solve():` — no input; it returns the two objects you need.
- `name = "distilbert-base-uncased"` — the base model id — plain DistilBERT, *not* fine-tuned for any task.
- `tokenizer = AutoTokenizer.from_pretrained(name)` — downloads the tokenizer that matches that model.
- `model = _____.from_pretrained(name, num_labels=_____)` — loads the model with a fresh classification head. Blanks = which Auto class to use and how many output labels.
- `return tokenizer, model` — hands both back — you need the pair to tokenize *and* to run/fine-tune the model.

Why this model. `distilbert-base-uncased` is the **plain**, *not*-fine-tuned DistilBERT (encoder-only). `AutoModelForSequenceClassification` bolts a **fresh, untrained** classification head on top — so its predictions are random until you fine-tune it. That's the whole point: this is the *starting point for training*, in contrast to ex5, which loaded the already-fine-tuned version.

How it classifies. DistilBERT turns the text into **contextual embeddings** — one vector per token. The vector of the special `[CLS]` token (the first one) is used as the **summary of the whole sentence**. The **classification head** (a small linear layer — in DistilBERT preceded by a pre-classifier + ReLU) maps that summary vector to one **logit** per class: raw, unbounded scores. To read them as probabilities you apply **softmax**; but to simply pick the predicted class you take the `argmax` of the logits directly — no softmax needed (that's what ex11 does). One caveat specific to this exercise: the head

is **freshly added and untrained**, so its logits are essentially random until you fine-tune the model (contrast with ex5, whose head is already trained).

Hint. Of the two imported classes, one adds a classification head — its name says so. How many classes does POSITIVE-vs-NEGATIVE have?

Complete this in `exercises/ex06_load_model.py`:

```
from transformers import AutoTokenizer, AutoModelForSequenceClassification

def solve():
    name = "distilbert-base-uncased"
    tokenizer = AutoTokenizer.from_pretrained(name)
    # Fill: which imported class loads a classification head, and how many labels.
    model = _____.from_pretrained(name, num_labels=____)
    return tokenizer, model
```

Exercise 7 · Tokenize texts (padding, truncation, tensors)

Task. Return padded, truncated (max 64) PyTorch tensors for a list of `texts`.

What the code does. The code calls the tokenizer on a list of texts. `padding=True` pads sequences to equal length so they form a rectangular batch; `truncation=True` cuts overly long ones. The blanks are the **tensor format** to return and the **max length** cap.

Line by line:

- `from transformers import AutoTokenizer` — imports the tokenizer class (used by the demo).
- `def solve(tokenizer, texts):` — receives a ready tokenizer and a list of texts.
- `return tokenizer(texts, return_tensors=____, padding=True, truncation=True, max_length=____)` — tokenizes the batch: `padding=True` pads all to equal length (a rectangular batch), `truncation=True` cuts long ones. Blanks = the tensor format and the length cap.

input_ids & attention_mask. The tokenizer returns a small dictionary of tensors:

- **input_ids** — the text turned into numbers: each token maps to an integer id from the model's vocabulary. For DistilBERT, 101 and 102 are the special [CLS] / [SEP] markers that wrap every sentence, and 0 is [PAD].
- **attention_mask** — a parallel array, same shape, with **1 = real token** and **0 = padding**. *How it protects the result:* inside the Transformer, attention makes each token look at the others through weights produced by a **softmax**. At every position where the mask is 0, the attention score is set to $-\infty$ **before the softmax**, and since $e^{(-\infty)} = 0$ those slots get weight **0** — so padding contributes nothing and real tokens attend only to real tokens. That's why you can pad as much as you like without changing the output.

- **Not the same as the causal mask** (§3): `attention_mask` says *which tokens are real*; the causal mask (decoders only) says *each token may see only earlier ones*. The model adds the causal one itself — the tokenizer gives you the padding one.
- You always pass both to the model together: `model(**enc)` (i.e. `model(input_ids=..., attention_mask=...)`).

Hint. PyTorch tensors have a two-letter code; the length cap is the number named in the task.

Complete this in `exercises/ex07_tokenize.py`:

```
from transformers import AutoTokenizer

def solve(tokenizer, texts):
    # Fill: return PyTorch tensors ("pt"), and cap the length at 64 tokens.
    return tokenizer(texts, return_tensors=____, padding=True,
                    truncation=True, max_length=____)
```

Exercise 8 · Tokenize a whole dataset with `.map`

Task. Tokenize a dataset row-by-row with `.map`, processing it in batches.

What the code does. `load_dataset("imdb", split="train[:1%]")` loads 1% of the train split; the inner `tok` function tokenizes the `text` column; `.map` applies it to every row. The blank controls whether `.map` passes rows **one at a time or in batches**.

The point. The goal of this exercise isn't to train the model — it's to **prepare the data**. The raw dataset held only `text` and its `label`; running it through `.map` with the tokenizer produces a new dataset that keeps those and adds `input_ids` and `attention_mask` — the format Transformers expect for fine-tuning. In effect you've turned a **human-readable** dataset into one a **neural network can process**: `input_ids` and `attention_mask` are the model's **inputs**, `label` is the **target** it must learn to predict, and `text` just rides along. (One technicality: right after `.map` the ids are still Python lists; the Trainer / data collator turns them into tensors at training time.)

Line by line:

- `from transformers import AutoTokenizer` — the tokenizer class (used by the demo).
- `from datasets import load_dataset` — the loader for ready-made datasets.
- `def solve(tokenizer):` — receives a tokenizer to apply.
- `data = load_dataset("imdb", split="train[:1%]")` — loads 1% of IMDB's train split — a small set of labelled movie reviews.
- `def tok(batch): ...` — a helper that tokenizes the `text` column of a batch of rows.
- `return data.map(tok, batched=____)` — applies `tok` across the whole dataset. Blank = process rows **in batches** (much faster).

Batched tokenization with `.map`. `data.map(tok, ...)` applies the tokenizer across the **whole dataset** instead of one sentence at a time, and returns a **new dataset** that keeps the original columns (`text`, `label`) and **adds** `input_ids` and `attention_mask` — the columns the model consumes during fine-tuning. Three details worth knowing:

- **Batch size:** processing *in batches* doesn't mean *all at once* — by default `tok` receives **1,000 rows** per call.
- **Where the speed comes from:** the big win is that **fast (Rust) tokenizers** process a whole batch in parallel; with a slow (pure-Python) tokenizer the gain is much smaller.
- **A padding gotcha:** `padding=True` here pads **within each batch**, so different batches can end up with different lengths. For training you often defer padding to a *data collator* (dynamic per-batch padding) or use `padding="max_length"` for a uniform size.

Hint. You want `.map` to process rows in batches rather than one at a time.

Complete this in `exercises/ex08_tokenize_dataset.py`:

```
from transformers import AutoTokenizer
from datasets import load_dataset

def solve(tokenizer):
    data = load_dataset("imdb", split="train[:1%]")

    def tok(batch):
        return tokenizer(batch["text"], padding=True, truncation=True, max_length=128)

    # Fill: apply tok across the dataset in batches.
    return data.map(tok, batched=_____)
```

5 · Fine-tuning through training

With a model and a tokenized dataset in hand, fine-tuning comes down to three objects.

TrainingArguments holds the hyperparameters — where to save (`output_dir`), when to evaluate (`eval_strategy="epoch"`), how long (`num_train_epochs`), how fast (`learning_rate`), the batch sizes, and regularization (`weight_decay`). **Trainer** ties everything together — the model, those args, the train and eval datasets, and the tokenizer — and `trainer.train()` runs the loop.

Once trained, using the model directly means doing by hand what the pipeline did for you: tokenize the text, run it under `torch.no_grad()` (no gradients needed for inference), read the raw **logits**, take the `argmax` for the predicted class id, and map that id to a label via `model.config.id2label`. Finally, `model.save_pretrained(dir)` and `tokenizer.save_pretrained(dir)` persist your work; `from_pretrained(dir)` loads it back.

Heads-up: in recent transformers, `TrainingArguments` and `Trainer` need the `accelerate` package (it's in `requirements.txt`). If you see "the Trainer ... requires accelerate", run `pip install -r requirements.txt` (or `pip install "accelerate>=1.1.0"`).

Exercises

Exercise 9 · Set up `TrainingArguments` · offline

Task. Return `TrainingArguments`: output `./out`, evaluate each epoch, 3 epochs, lr $2e-5$, train/eval batch 8, weight decay 0.01. *Offline*.

What the code does. This constructs the hyperparameter bundle for training. `output_dir` is where checkpoints go; each blank is one hyperparameter — when to evaluate, how many epochs, the learning rate, the train batch size, and weight decay — with its **target value written in the comment beside it**.

The point. This exercise doesn't train anything — it **defines the rules of training**. `TrainingArguments` is a configuration object the `Trainer` reads later to know how many passes to make over the data (**epochs**), how big each weight update is (**learning rate**), how many examples to process at once (**batch size**), how much to regularise (**weight decay**), when to **evaluate**, and where to **save checkpoints** (`output_dir`). It's the *flight plan* the fine-tuning run will follow — it has dozens of options, and here you set the essential ones.

Line by line:

- `from transformers import TrainingArguments` — imports the hyperparameter container.
- `def solve():` — returns the configuration object (no training happens here).
- `return TrainingArguments(output_dir="./out", eval_strategy=__, num_train_epochs=__, ...)` — builds the training config; `output_dir` is where checkpoints go and each blank is one hyperparameter (eval schedule, epochs, learning rate, train batch size, weight decay) — the target value is written in the comment beside each blank.

Key hyperparameters. Three settings are worth understanding, not just filling in:

- `num_train_epochs` (**epochs**) — one epoch is one full pass over the whole training set, so `3` means the model sees every example three times. Too few and it *underfits* (hasn't learned enough); too many and it starts to *memorise* the training data (**overfits**).
- `learning_rate` — the **size of each step** when the weights are updated. Big steps learn fast but can overshoot and diverge; tiny steps are stable but slow. $2e-5$

(0.00002) is a deliberately small value typical of fine-tuning: the pretrained weights are already good, so you only want to *nudge* them, not overwrite them.

- **weight_decay** — a **regularisation** term that, each step, gently pulls every weight toward zero (it penalises large weights). The penalty grows with the **square** of each weight (L2), so a few big weights cost far more than many small ones — e.g. two weights of 1 cost $1+1 = 2$, but one weight of 2 costs $2 \times 2 = 4$. The model therefore prefers to **spread** the signal across many features rather than lean hard on one, which **reduces overfitting**. It's a lever, though: too strong and the model is over-constrained and *underfits*. `0.01` is a mild, standard value. (Technical note: this is L2 regularisation; the `Trainer`'s AdamW optimiser applies it as *decoupled* weight decay — directly on the weights — but the intuition is the same.)

Hint. Each value you need is written in the comment next to its blank.

Complete this in `exercises/ex09_training_args.py`:

```
from transformers import TrainingArguments

def solve():
    # Fill the values described in the task.
    return TrainingArguments(
        output_dir="./out",
        eval_strategy=____,          # evaluate at the end of each epoch
        num_train_epochs=____,       # 3
        learning_rate=____,          # 2e-5
        per_device_train_batch_size=____, # 8
        per_device_eval_batch_size=8,
        weight_decay=____,           # 0.01
    )
```

Exercise 10 · Wire up the Trainer

Task. Return a `Trainer` built from the pieces (don't call `.train()`).

What the code does. In a single call you build the `Trainer` — the object that will run the fine-tuning loop — handing it the pieces from the previous exercises. `solve()` only *constructs* it; no training happens here.

The point. This is the moment **every piece you built connects into one pipeline**. There's no need to wire the arguments one by one to learn them — `model` (ex6), `TrainingArguments` (ex9), the train/eval datasets (ex8) and the `tokenizer` (ex6/7) were each covered already. The new idea is that the `Trainer` **orchestrates the whole fine-tuning**: hand it those pieces and it knows how to run the training loop.

Is the model being trained in this exercise? It depends what you run:

- `solve()` — the part you complete — only **builds** the `Trainer` object. **No training.**
- **Running the file** (`python exercises/ex10_build_trainer.py`) builds it *and* then calls `trainer.train()` on a tiny dataset for two steps — **yes, a real mini-train** you can watch (a loss is printed).

- `python check.py 10` only verifies the wiring. **No training.**

So the *exercise itself* just assembles the Trainer; the bundled *demo* adds a two-step train so the loop isn't abstract. And what `trainer.train()` actually does, each step, is: pull a **batch** from the `train_dataset` → run it **through the model** → compute the **loss** (how wrong it is) → **backpropagate** to get the gradients (which way each weight should move) → **update the weights** a little (step size = the learning rate) → **evaluate** when `eval_strategy` says to → **save checkpoints** to `output_dir`.

Hint. Each `Trainer` parameter takes the argument of the same name that `solve` received.

Complete this in `exercises/ex10_build_trainer.py`:

```
from transformers import (AutoTokenizer, AutoModelForSequenceClassification,
                          Trainer, TrainingArguments)
from datasets import Dataset

def solve(model, args, train_ds, eval_ds, tokenizer):
    # Fill: pass the model, the args, and the two datasets into the Trainer.
    return Trainer(
        model=____,
        args=____,
        train_dataset=____,
        eval_dataset=____,
        tokenizer=tokenizer,
    )
```

Exercise 11 · Use the fine-tuned model (from logits)

Task. Return `"POSITIVE"` / `"NEGATIVE"` for `text`, computed **from raw logits**.

The key idea. The goal isn't to learn `torch.argmax()` — it's to understand that a classification model **doesn't return a label**; it returns **logits** (raw numeric scores), and *you* have to turn those into the final class. This exercise does by hand what `pipeline()` was doing for you automatically.

Where we are. This is the last step of *inference*:

```
Text -> Tokenizer -> Model -> Logits -> pick the best class -> "POSITIVE"
```

What are logits? The model doesn't hand back `POSITIVE`; it returns a vector of scores, e.g.

```
tensor([[ -1.25,  3.82]])
```

one score per class (index `0` -> `NEGATIVE`, index `1` -> `POSITIVE`). These are **not** probabilities — they're raw scores. (You'd apply *softmax* to turn them into probabilities, but you don't need to just to pick the winner.)

How we get the prediction. `torch.argmax(logits, dim=1)` returns the *position* of the largest value. For `[-1.25, 3.82]` that is `1`, because `3.82` is the highest.

Why `dim=1` ? Logits have shape `[batch, classes]` — one row per sentence, one column per class. Reducing over `dim=1` (the class axis) picks the top class **for each sentence** in the batch.

How we get the label. The model carries a dictionary `model.config.id2label`, e.g. `{0: "NEGATIVE", 1: "POSITIVE"}`, so `model.config.id2label[pred]` turns the class id into a human-readable label.

Why `torch.no_grad()` ? At inference we don't want to compute gradients or update weights — we only read a prediction. Wrapping the forward pass in `with torch.no_grad():` makes it faster and use less memory.

Why the model matters. This is the same SST-2 fine-tuned DistilBERT as ex5. *Because* it's fine-tuned, its logits are meaningful and `argmax` gives a real POSITIVE/NEGATIVE; the untrained head from ex6 would give noise.

The full flow:

```
Text
|
v
Tokenizer -> input_ids
|
v
fine-tuned Model
|
v
Logits          e.g. [-1.25, 3.82]
|
v   argmax()    -> 1
id2label[1]
|
v
"POSITIVE"
```

What you really learn. This exercise reveals what was hidden inside `pipeline()` : a classification model returns a vector of **logits**, and making a prediction means selecting the highest-scoring class (`argmax`) and mapping its id to a label (`id2label`). That is what actually happens during inference with a Transformer.

Hint. Logits are shaped `[batch, classes]` , so reduce over the class axis; then match the capitalisation the task expects.

Complete this in `exercises/ex11_predict.py` :


```
import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification

def solve(text):
    name = "distilbert-base-uncased-finetuned-sst-2-english"
    tok = AutoTokenizer.from_pretrained(name)
    model = AutoModelForSequenceClassification.from_pretrained(name)
    inp = tok(text, return_tensors="pt", padding=True, truncation=True)
    with torch.no_grad():
        logits = model(**inp).logits
    # Fill: the axis to argmax over (1), and the string method to upper-case.
    pred = int(torch.argmax(logits, dim=____))
    return model.config.id2label[pred].____()
```

6 · Fine-tuning approaches & N-shot prompting

Full training isn't the only option. **Full fine-tuning** updates every weight — most powerful, most expensive. **Partial fine-tuning** freezes most layers and trains only a few — cheaper and less prone to overfitting. **Transfer learning** is the umbrella idea: take a model trained on one task and adapt it to a related one.

Often you don't need to train at all. **N-shot prompting** teaches the model *inside the prompt*, by example: **zero-shot** gives no example (just the instruction), **one-shot** gives a single worked example, and **few-shot** gives several. More examples generally sharpen the output format and accuracy — at the cost of a longer prompt.

7 · Where to look things up

Start here (official, free):

- **Hugging Face LLM/NLP Course** — the best structured, hands-on intro: <https://huggingface.co/learn>
- **Transformers docs (home)** — <https://huggingface.co/docs/transformers>

API reference for what you used in these exercises:

- `pipeline()` and the full task list — https://huggingface.co/docs/transformers/main_classes/pipelines
- **Auto classes** (`AutoTokenizer` , `AutoModelFor...`) — https://huggingface.co/docs/transformers/model_doc/auto
- **Tokenizer** (`padding` / `truncation` / `return_tensors`) — https://huggingface.co/docs/transformers/main_classes/tokenizer
- **Trainer & TrainingArguments** — https://huggingface.co/docs/transformers/main_classes/trainer
- **Datasets** (`load_dataset` , `.map`) — <https://huggingface.co/docs/datasets>

- **Task guides** (summarization, translation, classification, QA...) — <https://huggingface.co/docs/transformers/tasks/summarization>

Finding a model:

- **The Model Hub** — filter by task on the left: <https://huggingface.co/models>
- **Read the model card** — what architecture it is, its intended task and languages, and the exact `model=` id to copy.

Background / the "why":

- PyTorch docs — <https://pytorch.org/docs/stable/index.html>
- Papers: *Attention Is All You Need* (Transformer, 2017), *BERT* (2018), *GPT-2* (2019) — searchable on <https://arxiv.org>

A skill worth practising, not just bookmarking:

- In Python, `help(pipeline)` (or `TrainingArguments?` in a notebook) prints the docstring for any object — the fastest reference of all.
- Every argument you pass (`max_length`, `eval_strategy`, ...) has a one-line explanation in the class's docstring / API page — look it up there instead of guessing.

8 • Self-check

You can now: run a task with `pipeline()`; control generation length/padding; translate; tell encoder-only / decoder-only / encoder-decoder apart and verify in code; move from `pipeline` to Auto classes; tokenize (with padding/truncation) and `.map` a dataset; set up `TrainingArguments` + `Trainer`; run inference from logits; and prompt zero/one/few-shot. That's the whole arc of the two chapters.